

real-forge-signature-detection

April 15, 2024

1 Import Libraries

```
[1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import cv2
import os
from sklearn.model_selection import train_test_split
import tensorflow as tf
from tensorflow.keras.layers import Conv2D, MaxPooling2D, MaxPool2D, Dense, Flatten, Dropout
from tensorflow.keras.models import Sequential
from tensorflow.keras.callbacks import EarlyStopping, ReduceLROnPlateau
from tensorflow.keras.preprocessing.image import ImageDataGenerator
import visualkeras
```

2 Loading the Data and Normalizing it

```
[2]: # load the dataset
real_path = 'signature_dataset/original dataset'
forge_path = 'signature_dataset/fraud dataset'

# set the image size to 128x128
img_size = (128, 128)

real_images = []
for img_name in os.listdir(real_path):
    img = cv2.imread(os.path.join(real_path, img_name), cv2.COLOR_RGB2GRAY)
    img = cv2.resize(img, img_size)
    real_images.append(img)
real_images = np.array(real_images)

forge_images = []
for img_name in os.listdir(forge_path):
    img = cv2.imread(os.path.join(forge_path, img_name), cv2.COLOR_RGB2GRAY)
    img = cv2.resize(img, img_size)
```

```

    forge_images.append(img)
forge_images = np.array(forge_images)

# normalize the data
real_images = real_images.astype('float32') / 255.0
forge_images = forge_images.astype('float32') / 255.0

```

3 Create labels for the real and forged signatures

```

[3]: num_forge_images = len(forge_images)
    num_real_images = len(real_images)

    # Create labels for the real and forged signatures
    forge_labels = np.zeros(num_forge_images, dtype=int)
    real_labels = np.ones(num_real_images, dtype=int)

    # Concatenate the real and forged images and labels
    X = np.concatenate((real_images, forge_images), axis=0)
    y = np.concatenate((real_labels, forge_labels), axis=0)

```

```

[4]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
    ↪random_state=42, shuffle=True)

```

```

[5]: print(X_train.shape, X_test.shape, y_train.shape, y_test.shape)

(4891, 128, 128, 3) (1223, 128, 128, 3) (4891,) (1223,)

```

4 Making the CNN model

```

[6]: # Create a Sequential model
    model = Sequential()

    # Add a convolutional layer
    model.add(Conv2D(filters=16, kernel_size=(3,3), activation='relu',
    ↪input_shape=(128, 128, 3)))
    # Add a convolutional layer
    model.add(Conv2D(filters=16, kernel_size=(3,3), activation='relu'))
    # Add a max pooling layer
    model.add(MaxPooling2D(pool_size=(2,2)))
    # Add a dropout layer to reduce overfitting
    model.add(Dropout(rate=0.2))

    # Add a convolutional layer
    model.add(Conv2D(filters=32, kernel_size=(3,3), activation='relu'))

```

```

# Add a convolutional layer
model.add(Conv2D(filters=32, kernel_size=(3,3), activation='relu'))
# Add a max pooling layer
model.add(MaxPooling2D(pool_size=(2,2)))
# Add a dropout layer to reduce overfitting
model.add(Dropout(rate=0.2))

# Add a convolutional layer
model.add(Conv2D(filters=64, kernel_size=(3,3), activation='relu'))
# Add a convolutional layer
model.add(Conv2D(filters=64, kernel_size=(3,3), activation='relu'))
# Add a max pooling layer
model.add(MaxPooling2D(pool_size=(2,2)))
# Add a dropout layer to reduce overfitting
model.add(Dropout(rate=0.2))

# Flatten the output from the convolutional layers
model.add(Flatten())
# Add a fully connected layer with 128 neurons and a relu activation function
model.add(Dense(units=256, activation='relu'))
# Add a dropout layer to reduce overfitting
model.add(Dropout(rate=0.5))
# Add the output layer with a sigmoid activation function for binary
  ↪ classification
model.add(Dense(units=1, activation='sigmoid'))

# Print a summary of the model architecture
model.summary()

```

Model: "sequential"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 126, 126, 16)	448
conv2d_1 (Conv2D)	(None, 124, 124, 16)	2320
max_pooling2d (MaxPooling2D)	(None, 62, 62, 16)	0
dropout (Dropout)	(None, 62, 62, 16)	0
conv2d_2 (Conv2D)	(None, 60, 60, 32)	4640
conv2d_3 (Conv2D)	(None, 58, 58, 32)	9248

max_pooling2d_1 (MaxPooling2D)	(None, 29, 29, 32)	0
dropout_1 (Dropout)	(None, 29, 29, 32)	0
conv2d_4 (Conv2D)	(None, 27, 27, 64)	18496
conv2d_5 (Conv2D)	(None, 25, 25, 64)	36928
max_pooling2d_2 (MaxPooling2D)	(None, 12, 12, 64)	0
dropout_2 (Dropout)	(None, 12, 12, 64)	0
flatten (Flatten)	(None, 9216)	0
dense (Dense)	(None, 256)	2359552
dropout_3 (Dropout)	(None, 256)	0
dense_1 (Dense)	(None, 1)	257


```

=====
Total params: 2431889 (9.28 MB)
Trainable params: 2431889 (9.28 MB)
Non-trainable params: 0 (0.00 Byte)
-----

```

```
[7]: visualkeras.layered_view(model)
```

```
[7]:
```



```
[8]: # EarlyStopping callback
early_stopping = EarlyStopping(monitor='val_loss',           # Monitor validation_
                                ↪ loss
                                min_delta=0.001,           # Minimum change in_
                                ↪ the monitored quantity to qualify as an improvement
                                patience=4,                 # Number of epochs_
                                ↪ with no improvement after which training will be stopped
                                restore_best_weights=True,   # Restore model_
                                ↪ weights to the best iteration
                                verbose=0)                  # Verbosity mode (0:_
                                ↪ silent, 1: update messages)

# ReduceLROnPlateau callback
reduce_learning_rate = ReduceLROnPlateau(monitor='val_accuracy', # Monitor_
                                           ↪ validation accuracy
                                           patience=2,           # Number of_
                                           ↪ epochs with no improvement after which learning rate will be reduced
                                           factor=0.5,          # Factor by_
                                           ↪ which the learning rate will be reduced (new_lr = lr * factor)
                                           verbose=1)           # Verbosity_
                                           ↪ mode (0: silent, 1: update messages)
```

5 Evaluating the Model

```
[9]: model.compile(optimizer='adam',  
                  loss='binary_crossentropy',  
                  metrics=['accuracy'])
```

```
[10]: history = model.fit(  
    X_train, y_train,  
    validation_data=(X_test, y_test),  
    batch_size=32,  
    epochs=30,  
    callbacks=[early_stopping, reduce_learning_rate],  
    verbose=1  
)
```

```
Epoch 1/30  
98/98 [=====] - 68s 681ms/step - loss: 0.6913 -  
accuracy: 0.5557 - val_loss: 0.6342 - val_accuracy: 0.6329 - lr: 0.0010  
Epoch 2/30  
98/98 [=====] - 63s 641ms/step - loss: 0.6262 -  
accuracy: 0.6581 - val_loss: 0.5847 - val_accuracy: 0.6885 - lr: 0.0010  
Epoch 3/30  
98/98 [=====] - 75s 766ms/step - loss: 0.5930 -  
accuracy: 0.6866 - val_loss: 0.5541 - val_accuracy: 0.7244 - lr: 0.0010  
Epoch 4/30  
98/98 [=====] - 61s 621ms/step - loss: 0.5069 -  
accuracy: 0.7526 - val_loss: 0.4419 - val_accuracy: 0.7956 - lr: 0.0010  
Epoch 5/30  
98/98 [=====] - 64s 654ms/step - loss: 0.4137 -  
accuracy: 0.8141 - val_loss: 0.4053 - val_accuracy: 0.8144 - lr: 0.0010  
Epoch 6/30  
98/98 [=====] - 67s 681ms/step - loss: 0.3575 -  
accuracy: 0.8434 - val_loss: 0.3190 - val_accuracy: 0.8692 - lr: 0.0010  
Epoch 7/30  
98/98 [=====] - 59s 604ms/step - loss: 0.3076 -  
accuracy: 0.8669 - val_loss: 0.3093 - val_accuracy: 0.8823 - lr: 0.0010  
Epoch 8/30  
98/98 [=====] - 60s 615ms/step - loss: 0.2682 -  
accuracy: 0.8900 - val_loss: 0.2697 - val_accuracy: 0.8970 - lr: 0.0010  
Epoch 9/30  
98/98 [=====] - 64s 658ms/step - loss: 0.2316 -  
accuracy: 0.9053 - val_loss: 0.2627 - val_accuracy: 0.8962 - lr: 0.0010  
Epoch 10/30  
98/98 [=====] - 65s 661ms/step - loss: 0.1807 -  
accuracy: 0.9246 - val_loss: 0.2897 - val_accuracy: 0.8994 - lr: 0.0010  
Epoch 11/30  
98/98 [=====] - 63s 647ms/step - loss: 0.1490 -  
accuracy: 0.9442 - val_loss: 0.3017 - val_accuracy: 0.8782 - lr: 0.0010
```

```

Epoch 12/30
98/98 [=====] - 62s 630ms/step - loss: 0.1330 -
accuracy: 0.9448 - val_loss: 0.2507 - val_accuracy: 0.9092 - lr: 0.0010
Epoch 13/30
98/98 [=====] - 62s 636ms/step - loss: 0.1011 -
accuracy: 0.9605 - val_loss: 0.2496 - val_accuracy: 0.9158 - lr: 0.0010
Epoch 14/30
98/98 [=====] - 62s 633ms/step - loss: 0.1046 -
accuracy: 0.9644 - val_loss: 0.2157 - val_accuracy: 0.9231 - lr: 0.0010
Epoch 15/30
98/98 [=====] - 63s 643ms/step - loss: 0.0673 -
accuracy: 0.9753 - val_loss: 0.2993 - val_accuracy: 0.9150 - lr: 0.0010
Epoch 16/30
98/98 [=====] - ETA: 0s - loss: 0.0777 - accuracy:
0.9744
Epoch 16: ReduceLROnPlateau reducing learning rate to 0.0005000000237487257.
98/98 [=====] - 63s 641ms/step - loss: 0.0777 -
accuracy: 0.9744 - val_loss: 0.2354 - val_accuracy: 0.9231 - lr: 0.0010
Epoch 17/30
98/98 [=====] - 62s 634ms/step - loss: 0.0405 -
accuracy: 0.9865 - val_loss: 0.2601 - val_accuracy: 0.9305 - lr: 5.0000e-04
Epoch 18/30
98/98 [=====] - 66s 679ms/step - loss: 0.0324 -
accuracy: 0.9892 - val_loss: 0.2764 - val_accuracy: 0.9305 - lr: 5.0000e-04
Epoch 19/30
98/98 [=====] - ETA: 0s - loss: 0.0375 - accuracy:
0.9873
Epoch 19: ReduceLROnPlateau reducing learning rate to 0.0002500000118743628.
98/98 [=====] - 65s 660ms/step - loss: 0.0375 -
accuracy: 0.9873 - val_loss: 0.2734 - val_accuracy: 0.9289 - lr: 5.0000e-04

```

6 Testing Loss and Accuracy

```

[11]: train_loss, train_acc = model.evaluate(X_train, y_train, verbose=1)

print("The accuracy of the model for training data is:", train_acc * 100)
print("The Loss of the model for training data is:", train_loss)

```

```

153/153 [=====] - 9s 56ms/step - loss: 0.0308 -
accuracy: 0.9973
The accuracy of the model for training data is: 99.7342050075531
The Loss of the model for training data is: 0.03080846555531025

```

```

[12]: test_loss, test_acc = model.evaluate(X_test, y_test, verbose=1)

print("The accuracy of the model for testing data is:", test_acc * 100)
print("The Loss of the model for testing data is:", test_loss)

```

```
39/39 [=====] - 2s 54ms/step - loss: 0.2157 - accuracy: 0.9231
```

The accuracy of the model for testing data is: 92.31398105621338

The Loss of the model for testing data is: 0.21567381918430328

```
[13]: model.save("model/signature_detection.keras")
```

```
[16]: model.save("model/signature_detection.h5")
```

```
C:\Users\ASUS\AppData\Local\Programs\Python\Python310\lib\site-
packages\keras\src\engine\training.py:3079: UserWarning: You are saving your
model as an HDF5 file via `model.save()`. This file format is considered legacy.
We recommend using instead the native Keras format, e.g.
`model.save('my_model.keras')`.
  saving_api.save_model(
```

```
[17]: prediction = model.predict(X_test)
```

```
39/39 [=====] - 2s 52ms/step
```

```
[18]: prediction
```

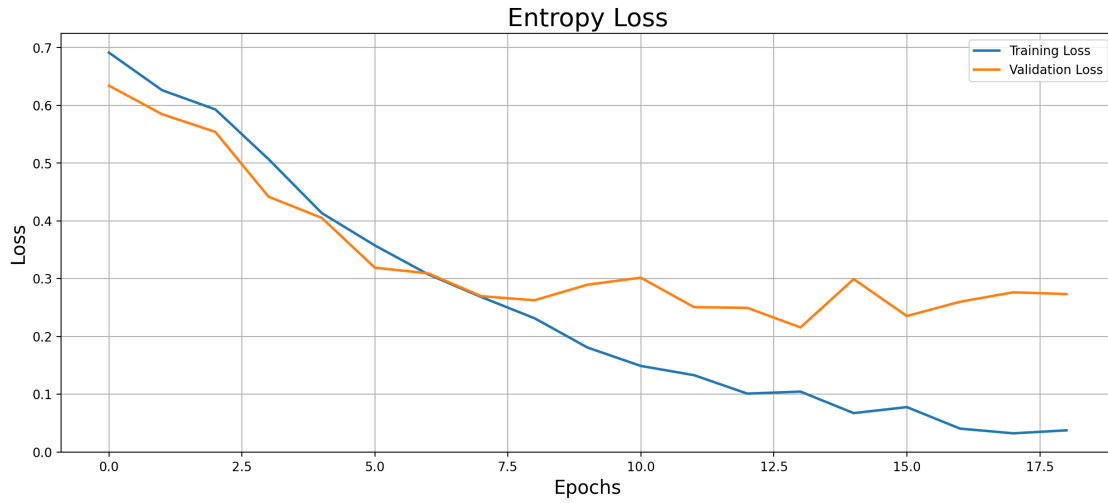
```
[18]: array([[0.00214481],
             [0.07519121],
             [0.02211507],
             ...,
             [0.02339804],
             [0.00028699],
             [0.00366422]], dtype=float32)
```

7 Plotting

```
[19]: plt.figure(figsize=(15, 6), dpi=200)

plt.title('Entropy Loss', fontsize=20)
plt.xlabel('Epochs', fontsize=15)
plt.ylabel('Loss', fontsize=15)
plt.plot(history.history['loss'], label='Training Loss', linewidth=2)
plt.plot(history.history['val_loss'], label='Validation Loss', linewidth=2)
plt.legend()
plt.grid(True)

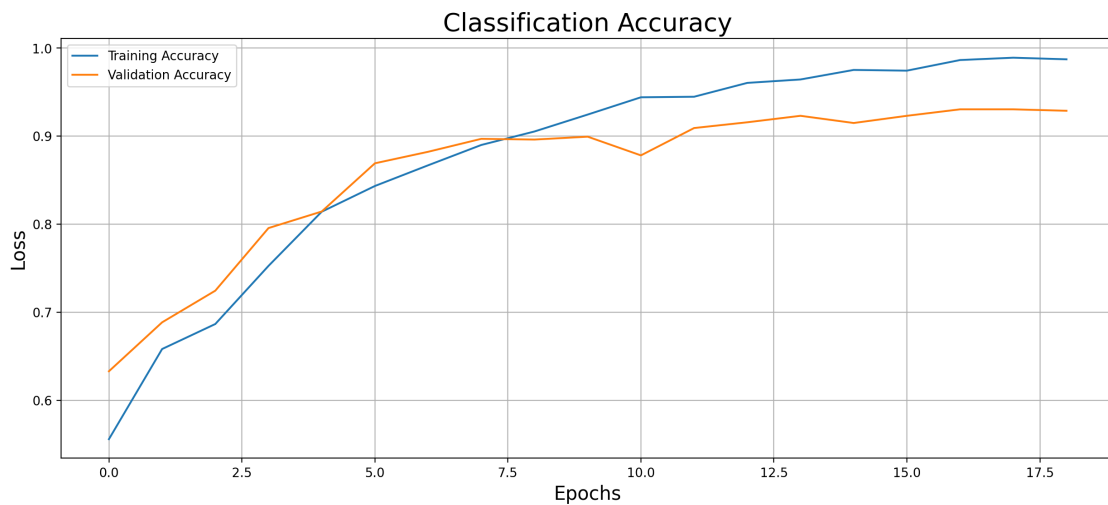
plt.show()
```

```
[20]: plt.figure(figsize=(15, 6), dpi=200)

plt.title('Classification Accuracy', fontsize=20)
plt.xlabel('Epochs', fontsize=15)
plt.ylabel('Loss', fontsize=15)
plt.plot(history.history['accuracy'], label='Training Accuracy')
plt.plot(history.history['val_accuracy'], label = 'Validation Accuracy')
plt.legend()
plt.grid(True)

plt.show()
```



8 Detection of Real and Forged Signature

```
[27]: # Load a signature image
# You can change the image path and check if it is forged or real
# img = cv2.imread('dataset/train/original dataset/original_01_01.jpg', cv2.
      ↪IMREAD_GRAYSCALE)
img = cv2.imread('signature_dataset/fraud dataset/forgeries_12_23.png', cv2.
      ↪COLOR_RGB2GRAY)
img = cv2.resize(img, (128, 128))
img = np.array(img).reshape(1, 128, 128, 3) / 255.0

# Predict the class of the signature image
prediction = model.predict(img)

if prediction > 0.7:
    print("The signature is real.")
else:
    print("The signature is forged.")
```

1/1 [=====] - 0s 42ms/step

The signature is forged.

```
[28]: # Load a signature image
# You can change the image path and check if it is forged or real
img = cv2.imread('signature_dataset/original dataset/original_13_20.png', cv2.
      ↪COLOR_RGB2GRAY)
# img = cv2.imread('dataset/train/fraud dataset/forgeries_12_23.png', cv2.
      ↪IMREAD_GRAYSCALE)
img = cv2.resize(img, (128, 128))
img = np.array(img).reshape(1, 128, 128, 3) / 255.0

# Predict the class of the signature image
prediction = model.predict(img)

if prediction > 0.7:
    print("The signature is real.")
else:
    print("The signature is forged.")
```

1/1 [=====] - 0s 43ms/step

The signature is real.

```
[ ]:
```